

Amazon Web Services

DynamoDB, Redshift, Kinesis and Lambda

Markus Dale, November 2015

Managed Services With Rich Environment

- ▶ EC2 - Elastic Load Balancer service, Auto-Scaling, Security groups...
- ▶ S3 - Server-Side Encryption, Versioning, Notification (Object created/removed)...
- ▶ RDS - Relational Database Service - managed RDBMS
 - ▶ Supports Aurora, MySQL, PostgreSQL, Maria, Oracle, SQL Server
 - ▶ Multi-AZ, read replicas, backups

Developer productivity - mission/business over infrastructure

Right Tool for the Right Job

- ▶ When you need that corkscrew...
- ▶ DynamoDB - NoSQL database
- ▶ Redshift - data warehouse
- ▶ Kinesis - Streaming lots of data
- ▶ Lambda - run code on event or timer

AWS DynamoDB

- ▶ Managed NoSQL database since 2012
 - ▶ Dynamo key-value paper 2007 (open source Cassandra, Riak...)
- ▶ Right tool: High write throughput, query on few attributes, evolving schema, small rows
- ▶ Cost: Provisioned read/write throughput and actual storage
- ▶ Scalable (up and down via API)
- ▶ fault-tolerant: synchronously replicate across multi-AZ
- ▶ SSD storage
- ▶ No updates/patching

DynamoDB Tables

- ▶ Tables have primary key with optional range key (sort key)
- ▶ Use primary key to determine partition (should have high variability!)
- ▶ Can query on primary key (and range key)
- ▶ Each item (~ row) has unique primary key + range key, attribute (key/value)
- ▶ No fixed schema (other than primary key)
- ▶ 400KB per item (store reference to BLOBs in S3)
- ▶ Primary key: String, Number, or Binary
- ▶ Attributes
 - ▶ Number, String, Binary, Boolean, and Null.
 - ▶ Document types – List and Map.
 - ▶ Set types – String Set, Number Set, and Binary Set

DynamoDB Provisioned Capacity

- ▶ Read capacity: Number of item reads per second \times 4 KB item size (x2 for eventually consistent)
- ▶ Write capacity: Number of item writes per second \times 1 KB item size

DynamoDB Sample Table - Reply

- ▶ Hash Key: Amazon DynamoDB#DynamoDB Thread 1
(denormalized - reference Forum table)
- ▶ Range Key: ReplyDateTime
- ▶ Scan (with filter)
- ▶ Query - by hash (with optional range)

DynamoDB Query

- ▶ On Hash Key
- ▶ On Hash Key, Range Key
- ▶ Can add filter
- ▶ Secondary Index - local/global

Amazon Redshift

- ▶ Fully managed, petabyte-scale data warehouse
- ▶ Online analytic processing (OLAP) and business intelligence (BI) - complex queries at scale
 - ▶ Joins!
- ▶ Use SQL (PostgreSQL 8.0.2 syntax)
- ▶ Single or multi-node cluster in a single AZ
- ▶ Continuously backed up to S3
- ▶ Can also enable manual/automated snapshots
- ▶ \$1,000 per terabyte per year (often 3x compression - \$333/TB/year)

Redshift Architecture

- ▶ Client (JDBC/ODBC) to Leader node
- ▶ One to many compute nodes (same AZ)
- ▶ Massively parallel processing (MPP) - query optimizer
- ▶ Columnar data storage

Redshift Integration

- ▶ Load/store in S3
- ▶ DynamoDB
- ▶ Elastic MapReduce
- ▶ Kinesis Streams

AWS Kinesis Streams

- ▶ Managed real-time data processing (~ Apache Kafka)
- ▶ Stores data up to 7 days for each stream
- ▶ Reliable, durable rolling buffer - replicates across 3 AZs
- ▶ Can have multiple Kinesis applications process same stream
 - ▶ Real-time application monitoring, logs
 - ▶ Feed dashboards
 - ▶ Alerts via SNS
 - ▶ IoT

Kinesis Streams Scaling and Data Records

- ▶ Add shards: 1 shard = 1 MB/second write, 2MB/second reads, 1000 puts/second
- ▶ From MB to TB, 1000s to millions of PUT records per second
- ▶ Sequence Number (assigned by Kinesis)
- ▶ Partition Key (user assigned)
- ▶ Data blob (after base64 encoding $< 1\text{MB}$)

AWS Lambda

- ▶ Event-driven compute
- ▶ Execute without provisioning/managing server
- ▶ Managed execution and scaling
- ▶ Python, Java, Node.js (JavaScript) - can call libraries, other languages

AWS Lambda Triggers Examples

- ▶ DynamoDB Table Update
- ▶ S3 Object modifications
- ▶ Amazon CloudWatch log entry
- ▶ Simple Email Service incoming email
- ▶ Kinesis Streams message
- ▶ Cron schedule

Lambda Execution

- ▶ Stateless (store state in S3 or DynamoDB)
- ▶ Charged by 100ms increments, number of requests, GB of RAM allocated
- ▶ Default: 3sec execution limit (max 300sec)
- ▶ Default: 100 simultaneous executions (can increase)
- ▶ IAM policies/roles to manage permissions (Identification & Access Management)