

What's in Apache Spark 2.0.0?

- ▶ Dataset evolution: SparkSession (entry to Dataset/DataFrame - not SQLContext(sc))
- ▶ Off-heap caching: Overcome GC limits, compressed object pointers (compressed oops - up to 32GB Java heap - 32 bits/4 bytes, 64 bits/8 bytes if over)

Project Tungsten - Closer to bare metal

- ▶ Memory management and binary processing
 - ▶ Java serialized object: JVM GC, 2 bytes UTF-16 encoding, header, hash code
 - ▶ C-style memory access - `sun.misc.Unsafe`
 - ▶ `allocateMemory`, `copyMemory`, `freeMemory`, `getAddress`, `getInt`, `putInt`
 - ▶ Spark manages memory
- ▶ Code generation: Don't create object, compare binary

Catalyst Optimizer

- ▶ represent query as tree/manipulate
- ▶ rule-based and cost-based optimization
- ▶ analysis, logical optimization, physical planning, and code generation to compile parts of queries to Java bytecode
- ▶ `Literal(value: Int)`, `Attribute(name: String)`
- ▶ `Add(left: TreeNode, right: TreeNode)`
- ▶ Rule: for example `tree.transform (add(lit1,lit2) = lit(1+2))`
- ▶ Analysis: look up column names/types/tables from Catalog
- ▶ Logical Optimization: rule-based with constant folding, predicate pushdown, projection pruning, null propagation, Boolean expression simplification, and other rules
- ▶ Physical plans - cost model, code gen- expressions (+ predicate pushdown)

Project Tungsten 2.0 - reduce CPU bottlenecks

- ▶ Virtual function calls
- ▶ Use CPU registers instead of cache/memory

Simple aggregate query with filter

- ▶ count how many items have the sk 1000

Pre-2.0 Apache Spark: Volcano Iterator Model

- ▶ Graefe, 1994 paper "Volcano" - iterator, virtual function call
- ▶ "elegantly compose arbitrary combinations of operators"